

# Two Methods

## 1. Kubernetes Setup using Kubeadm

**~Start - Execute the below commands in both Master/worker nodes**

**Login to both instances execute the below commands:**

```
sudo apt-get update -y && sudo apt-get install apt-transport-https -y
```

**Change to root user**

```
sudo su -
```

```
sudo curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key add -
```

```
cat <<EOF >/etc/apt/sources.list.d/kubernetes.list
deb https://apt.kubernetes.io/ kubernetes-xenial main
EOF
```

```
sudo apt-get update
```

**#Disable swap memory for better performance**

```
swapoff -a
```

```
sudo sed -i ' / swap / s/^(.*)$/#\1/g' /etc/fstab
```

**Enable IP tables**

#We need to enable IP tables for pod to pod communication.

```
modprobe br_netfilter
```

```
sysctl -p
```

```
sudo sysctl net.bridge.bridge-nf-call-iptables=1
```

**Install Docker on both Master and Worker nodes**

```
apt-get install docker.io -y
```

**Add ubuntu user to Docker group**

```
usermod -aG docker ubuntu
```

```
systemctl restart docker
```

```
systemctl enable docker.service
```

Type exit to come out of root user.

**Install Kubernetes Modules**

```
sudo apt-get install -y kubelet kubeadm kubectl kubernetes-cni
```

```
sudo systemctl daemon-reload
```

```
sudo systemctl start kubelet
```

```
sudo systemctl enable kubelet.service
```

```
sudo systemctl status docker
```

**#End - Execute the above commands in both Master/worker nodes#####**

```
cd /etc/docker/  
vi daemon.json
```

add this below commands:-

```
{  
  "exec-opts": ["native.cgroupdriver=systemd"]  
}
```

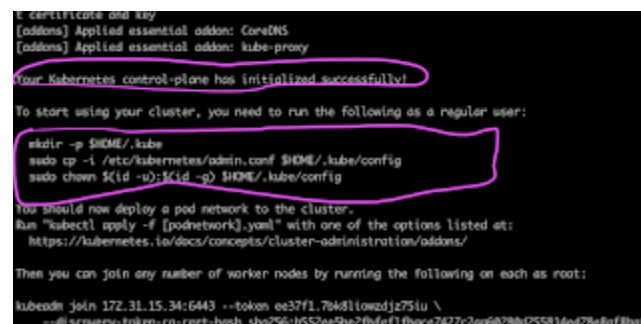
```
sudo systemctl daemon-reload  
sudo systemctl restart docker  
sudo systemctl restart kubelet
```

**Wait for Sometime, It will take some time**

### **Initialize Kubeadm on Master Node(only on Master Node)**

#Execute the below command as root user to initialize Kubernetes Master node.

```
sudo su -  
kubeadm init
```



```
...  
[addons] Applied essential addon: CoreDNS  
[addons] Applied essential addon: kube-proxy  
  
Your Kubernetes control-plane has initialized successfully!  
  
To start using your cluster, you need to run the following as a regular user:  
  
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config  
  
You should now deploy a pod network to the cluster.  
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:  
https://kubernetes.io/docs/concepts/cluster-administration/addons/  
  
Then you can join any number of worker nodes by running the following on each as root:  
  
kubeadm join 172.31.15.34:6443 --token ee37f1.7b48liowadjz75iu \
```

Make sure you see the above message to confirm master node is up.

**#Now type exit to exit from root user and execute below commands as normal user**

```
mkdir -p $HOME/.kube  
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config  
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

## Installing the Weave Net Add-On

`kubectl apply -f https://github.com/weaveworks/weave/releases/download/v2.8.1/weave-daemonset-k8s.yaml`

It may take a few mins to execute the above command and show the below message.

```
ubuntu@ip-172-31-90-196:~$ kubectl apply -f https://github.com/weaveworks/weave/releases/download/v2.8.1/weave-daemonset-k8s.yaml
serviceaccount/weave-net created
clusterrole.rbac.authorization.k8s.io/weave-net created
clusterrolebinding.rbac.authorization.k8s.io/weave-net created
role.rbac.authorization.k8s.io/weave-net created
rolebinding.rbac.authorization.k8s.io/weave-net created
daemonset.apps/weave-net created
```

Now execute the below command to see the pods.

`kubectl get pods --all-namespaces`

```
ubuntu@ip-172-31-28-68:~$ kubectl get pods --all-namespaces
NAMESPACE   NAME                                     READY   STATUS    RESTARTS   AGE
kube-system  coredns-66bff467f8-p452k               1/1     Running   0           158m
kube-system  coredns-66bff467f8-s2pd5               1/1     Running   0           158m
kube-system  etcd-ip-172-31-28-68                   1/1     Running   0           158m
kube-system  kube-apiserver-ip-172-31-28-68         1/1     Running   0           158m
kube-system  kube-controller-manager-ip-172-31-28-68 1/1     Running   0           158m
kube-system  kube-proxy-rg3dq                       1/1     Running   0           155m
kube-system  kube-proxy-w6r62                       1/1     Running   0           158m
kube-system  kube-scheduler-ip-172-31-28-68         1/1     Running   0           158m
kube-system  weave-net-f7b2j                        2/2     Running   1           155m
kube-system  weave-net-ftmod                        2/2     Running   0           158m
ubuntu@ip-172-31-28-68:~$
```

## Now login to Worker Node

### Join worker node to Master Node

The below command will join worker node to master node, execute this as a normal user by putting `sudo` before:

```
sudo kubeadm join <master_node_ip>:6443 --token xrvked.s0n9771cd9x8a9oc \
--discovery-token-ca-cert-hash sha256:288084720b5aad132787665cb73b9c530763cd1cba10e12574b4e97452137b4a
```

```
ubuntu@ip-172-31-21-242:~$ sudo kubeadm join 172.31.28.68:6443 --token npluq.tsz2ad8d42f3ar3x \
--discovery-token-ca-cert-hash sha256:3ba58d285788cf93178bc2a3d74dc5942896cf4775d9ded7d0ba
e53976
[INFO] 2020-09-24 10:40:40 [980] join.go:346 [preflight] WARNING: JoinControlPlane.controlPlane settings
will be ignored when control-plane flag is not set.
[preflight] Running pre-flight checks
[WARNING IsKubernetesCheck]: detected "cgroupfs" as the Docker cgroup driver. The recommen
ed driver is "systemd". Please follow the guide at https://kubernetes.io/docs/setup/cni/
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get kubeconfig
.yaml'
[kubelet-start] Downloading configuration for the kubelet from the "kubelet-config-1.18" ConfigMap in
the kube-system namespace
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.e
nvironment"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS bootstrap...

This node has joined the cluster:
* Certificate signing request was sent to apiserver and a response was received.
* The kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.
```

## Go to Master and type the below command

`kubectl get nodes`

the above command should display both Master and worker nodes.

```
ubuntu@ip-172-31-28-60:~$ kubectl get nodes
NAME                STATUS    ROLES    AGE   VERSION
ip-172-31-21-242    Ready    <none>   146m  v1.18.3
ip-172-31-28-60     Ready    master   150m  v1.18.3
ubuntu@ip-172-31-28-60:~$
```

It means Kubernetes Cluster - both Master and worker nodes are setup successfully and up and running!!!

(OR)

## 2. How to set up Kubernetes master-slave architecture?



Today, we will set up a complete **Kubernetes master-slave architecture using kubeadm**. According to the kubeadm source, Kubeadm is a tool built to provide kubeadm init and kubeadm join as best-practice “fast paths” for creating Kubernetes clusters.

We will consider building a Kubernetes setup with one master node and 2 worker nodes.

**Here we go,**

Let us assume we have three Ubuntu Linux machines named kmaster and knode

## 1. Installing Docker as the container runtime Interface

On all the machines do the following:

```
#update the repository
sudo apt-get update

#Install docker
sudo apt install docker.io

#Start and automate docker to start at run time
sudo systemctl start docker
sudo systemctl enable docker

#verify docker installation
docker container ls
```

**Kubeadm** will by default use **docker** as the container runtime interface. In case a machine has both docker and other container runtimes like **contained**, docker takes precedence.

## 2. Installing kubeadm tool

```
#add the required repository for kubeadm
sudo apt-get update && sudo apt-get install -y apt-transport-https curl
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo apt-
```

```
key add -
cat <<EOF | sudo tee /etc/apt/sources.list.d/kubernetes.list
deb https://apt.kubernetes.io/ kubernetes-xenial main
EOF

#update the repository
$ sudo apt-get update

#installing kubelet, kubeadm and kubectl
sudo apt-get install -y kubelet kubeadm kubectl

#setting apt-mark
sudo apt-mark hold kubelet kubeadm kubectl
```

apt-mark will change whether a package has been marked as being automatically installed. Hold is used to mark a package as held back, which will prevent the package from being automatically installed, upgraded, or removed.

#### Restart the kubelet if required

```
systemctl daemon-reload
systemctl restart kubelet
```

### 3. Initializing the control plane or making the node as master(on master node)

kubeadm init will initialize this machine to make it a master.

Kubernetes assigns each node a range of IP addresses, a CIDR(Classless Inter-Domain Routing) block so that each Pod can have a unique IP address. We will specify the private CIDR for the pods to be created.

```
kubeadm init --apiserver-advertise-address=192.168.56.101 --pod-network-cidr=192.168.0.0/16
```

```
Your Kubernetes master has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

You can now join any number of machines by running the following on each node
as root:

kubeadm join 192.168.56.101:6443 --token rzwx3l.dax3fxamhkegtz0w --discovery-token-ca-cert-hash sha256:dd78bef8526bdaebbda8da547c6bf387648dcadd65ae76dccc0defa1a5800f2
```

Now as seen in the output above, we need to run the below commands as a normal user to use the kubectl from terminal.

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Now the machine is initialized as master.

## 4. Joining Cluster

To join Kubernetes cluster simply copy that token at the end and paste it in worker node's terminal in super user mode.

```
#something like this
kubeadm join <control-plane-host>:<control-plane-port> --token <token> --
discovery-token-ca-cert-hash sha256:<hash>
```

That's it!

To see all running pods,

```
kubectl get pods -o wide --all-namespaces
```

You would see all pods are running except DNS one. To resolve that problem, enable **Calico network**.

```
# on master node
kubectl apply -f
https://docs.projectcalico.org/v3.14/manifests/calico.yaml
```

### **Elaboration on Calico related issue:**

You must deploy a Container Network Interface (CNI) based Pod network add-on so that your Pods can communicate with each other. Cluster DNS (CoreDNS) will not start up before a network is installed. We will use Calico as our CNI tool. Calico is a networking and network policy provider. Calico supports a flexible set of networking options so you can choose the most efficient option for your situation.